

1 Two Changes to the Model

Two changes relative to the model section of the circulated draft matter for the calibration.

First, I drop the Stone–Geary structure of the draft, in which children entered through the committed consumption and housing bundles $(\bar{c}(n, s), \bar{h}(n, s))$, and model family size with an equivalence scale, the more common device in the quantitative lifecycle literature. An equivalence scale measures how much more a family must spend than a childless couple to reach the same living standard. Consumption is then evaluated per adult equivalent, so a given bundle delivers less utility, and a marginal dollar more utility, in a larger household. [Fernández-Villaverde and Krueger \(2007\)](#) attribute about half of the lifecycle consumption hump to household size measured this way, and [Scholz et al. \(2006\)](#) and [Borella et al. \(2023\)](#) use scales of this form in lifecycle models close to this one. Within-period utility for a household with n dependent children at home is

$$u(c, h, n) = \frac{(c^{\alpha(n)} h^{1-\alpha(n)} / e(n))^{1-\sigma}}{1-\sigma} + \psi n, \quad \alpha(n) = \alpha_0 - (\delta_{\text{jump}} \mathbf{1}\{n \geq 1\} + \delta_a n), \quad (1)$$

Children raise the marginal value of the consumption–housing composite through $e(n)$ ¹ and tilt spending toward housing through $\alpha(n)$, both only while they live at home; once children leave, preferences revert to the homothetic case, and ψ is the direct utility flow per child at home. There are no minimum-consumption or minimum-housing requirements, so no household can be pushed below a subsistence bundle by a bad income draw.

Second, fertility is sequential rather than one-shot. Each period during the fertile ages, a childless household chooses between waiting and trying for a first birth, and a household with $n < 3$ children and a child still at home chooses between stopping and trying for child $n + 1$. An attempt at age a succeeds with probability π_a , which declines with age and reaches zero at 45, as in [Sommer \(2016\)](#).² No household chooses a number of children: completed family size is what the sequence of attempts and biology deliver. Both decisions carry type-1 extreme-value taste shocks, with separate scales κ_E on entry and κ_C on continuation, so that the extensive and intensive margins of fertility are governed by distinct dispersion parameters.

2 Calibration

I organize the model parameters into two groups. The first consists of parameters assigned from external evidence, measured directly in the data, or fixed as maintained restrictions. The second consists of ten parameters estimated internally by matching twelve model moments to their empirical counterparts. I describe the two groups in turn, explain which moments discipline which parameters, and then report the target system.

2.1 Parameters Set Outside the Estimation

I first describe the parameters assigned outside the estimation. Table 1 summarizes their values and sources.

¹ $e(n) = ((2 + 0.7n)/2)^{0.7}$, the scale of Citro and Michael (1995) used by [Scholz et al. \(2006\)](#), normalized to one for a childless couple.

²The failure probability rises exponentially in age, $1 - \pi_a = \omega_1 e^{\omega_2(a-18)}$ with $\omega_1 = 0.02$ and $\omega_2 = 0.134$, and $\pi_a = 0$ from age 45. Because the model period is four years, π_a has a direct empirical counterpart in the share of women who conceive within four years of starting: [Leridon \(2004\)](#) reports 0.91, 0.84, and 0.64 at ages 30, 35, and 40 under natural conditions, against 0.90, 0.80, and 0.62 here.

Demographics and preferences. A model period is four years. Households enter adult life at age 18, retire at age 66, and live at most through age 85. Survival is certain before retirement and thereafter follows four-year probabilities constructed from the combined male and female survivor counts in the 2023 Social Security period life table. Dependent children remain in the household for an expected 18 years, implying a maturity probability of $2/9$ per period. The number of children n takes the values 0, 1, 2, and 3 or more. I set risk aversion to $\sigma = 2$. Without children at home a household spends a share $1 - \alpha_0$ of its expenditure on housing, so α_0 is measured directly: I set $\alpha_0 = 0.733$, so that housing takes 26.7 percent of expenditure, its share among childless cash renters in the 2019–2023 Consumer Expenditure Survey.

Housing finance and transaction costs. Following Greaney et al. (2025), I set the annual real return on liquid wealth to $i_{\text{ann}} = 0.02$, annual housing depreciation to $\delta_{\text{ann}} = 0.011$, the proportional selling cost to $\psi^s = 0.06$, and the labor-income tax rate to $\tau_{\text{pay}} = 0.179$. The return is based on the average ten-year real interest rate, depreciation combines BEA residential-structure depreciation with an adjustment for the land share of house values, and the tax rate comes from OECD data. I set the annual property-tax rate to $\tau_{\text{ann}}^p = 0.01$. Mortgage debt cannot exceed a share $\phi = 0.80$ of the value of the house, so a purchase requires at least 20 percent equity. The benchmark imposes no separate payment-to-income limit.

Income process. I impose a deterministic lifecycle profile of household earnings, ε_a , normalized to average one over working ages. Within age, log income follows an annual AR(1) process, which I sample at the model’s four-year frequency and approximate with a five-state Rouwenhorst chain. I set the persistence to $\rho_z = 0.9136$ and the annual innovation standard deviation to $\sigma_z = 0.169$. Persistence and gross risk are the PSID estimates of Flodén and Lindé (2001), net of the survey measurement error they identify separately. Households bear this risk after taxes, so I scale the innovation by $1 - \tau$, using the progressivity $\tau = 0.181$ that Heathcote et al. (2017) estimate for the United States.

Housing products and supply. Housing is indivisible and available in a limited assortment of types and sizes (Davis and Van Nieuwerburgh, 2015). I represent this lumpiness with five owner-occupied house sizes, $\{H_k\}_{k=1}^5 = \{2, 4, 6, 8, 10\}$ rooms. Renters choose housing continuously up to $\bar{h}^R = 6$ rooms, the 90th percentile of the rental-room distribution in the matched ACS sample. These menus summarize the observed tenure-specific distribution of rooms. I set the housing-supply elasticity to $\eta = 1.75$, following the population-weighted metropolitan estimate in Saiz (2010).

Entry wealth. Households enter at age 18 with a draw from the PSID net-worth-to-income distribution of childless renters aged 18–24.

Table 1: Parameters set outside the estimation

Parameter	Value	Source or restriction
Model period	4 years	Model timing
Entry, retirement, maximum ages	18, 66, 85	Model timing
Post-retirement survival	Age-specific	2023 Social Security life table
Years with dependent children	18	Child-maturity restriction
Risk aversion σ	2	
Equivalence scale $e(n)$	$((2 + 0.7n)/2)^{0.7}$	Scholz et al. (2006)
Consumption share α_0	0.733	CEX childless cash renters, 2019–2023
Fecundity by age π_a	Declining; zero from 45	Leridon (2004)
Income persistence ρ_z	0.9136	Flodén and Lindé (2001)
Income innovation s.d. σ_z	0.169	Flodén and Lindé (2001), scaled by $1 - \tau$
Tax progressivity τ	0.181	Heathcote et al. (2017)
Annual real return i_{ann}	2.0%	Greaney et al. (2025)
Housing depreciation δ_{ann}	1.1%	Greaney et al. (2025)
Property tax τ_{ann}^p	1.0%	Maintained tax rate
Financed share ϕ	0.80	Davis and Van Nieuwerburgh (2015)
Sale cost ψ^s	6%	Davis and Van Nieuwerburgh (2015)
Labor-income tax τ_{pay}	17.9%	OECD U.S. average
Owner house sizes $\{H_k\}$	$\{2, 4, 6, 8, 10\}$ rooms	Matched ACS distributions
Largest rental unit $\bar{h}^R$	6 rooms	Matched ACS distributions
Housing-supply elasticity η	1.75	Saiz (2010)
Entrant wealth distribution	Wealth-to-income draws	PSID childless renters 18–24

2.2 Internally Estimated Parameters

I estimate the remaining ten parameters internally against twelve moments. Table 2 reports the current estimates. The parameters are chosen jointly. I group them into four economic blocks and, for each, describe which moments are most informative.

Table 2: Internally estimated parameters (provisional and illustrative)

Parameter	Interpretation	Estimate
β_{annual}	Annual discount factor	0.9821
δ_{jump}	Expenditure-share shift at entry into parenthood	0.0316
δ_a	Expenditure-share shift per additional child	0.0177
ψ	Utility flow per child at home	2.9968
κ_E	Taste-shock scale, first-birth entry decision	36.34
κ_C	Taste-shock scale, continuation decisions	0.2213
χ_O	Utility premium from owner-occupied housing	1.0276
$\bar{H}$	Housing-supply scale	8.8470
θ_0	Strength of the bequest motive	0.1117
θ_1	Bequest curvature shift	0.0465

Fertility. The utility value of children ψ raises the return to every birth, so it governs how many children households have. Completed fertility, the average number of children a woman has borne by the end of her fertile years, is the moment that most closely identifies it. The entry noise κ_E spreads the age at which women start trying, and biology turns late starts into late or missing births, so it

governs the timing of first births; the mean age at first birth and the share of first births after 30 are the informative moments here. The continuation noise κ_C governs how readily a household that has started keeps going: when it is small, starting a family is close to a commitment to several children, so fewer women start at all. Childlessness is mostly informed by this margin. The three are estimated together, because the same fertility outcomes respond to all of them, and the progression rates between birth orders are left untargeted. This assignment is the intended one; the estimates below do not yet deliver it, and I return to that in the fit.

Family housing. The two shifters δ_{jump} and δ_a set how far a household tilts spending toward housing when children arrive. The rooms a household adds at the first child identify the total tilt at entry into parenthood, $\delta_{\text{jump}} + \delta_a$, since a first-time parent takes both the jump and one per-child step; the rooms difference between families with three or more children and families with one or two pins down δ_a alone, since the jump is common to both and drops out. The family ownership gap has no shifter of its own. It informs both tilts, and to a lesser extent who remains childless, so it enters as an overidentifying restriction.

Tenure and supply. The owner premium χ_O raises the utility of owning over renting, and the prime-age ownership rate is the moment that most closely identifies it. The supply scale $\bar{H}$ sets the level of the housing stock; aggregate occupied rooms is most informative about it.

Saving and bequests. The discount factor β_{annual} governs how much wealth households carry through life, and aggregate net worth relative to annual gross labor earnings, 6.873 in the PSID, is the moment that most closely identifies it. The bequest strength θ_0 governs how much households hold back to leave at death; it is mostly informed by the annual flow of bequests relative to aggregate wealth, 0.0088 in [De Nardi and Yang \(2014\)](#). The bequest curvature θ_1 governs how the motive strengthens with wealth, and the informative moment is the ratio of the 90th to the 50th percentile of total wealth late in life, 3.448. Without the flow moment nothing in the estimation disciplines θ_0 , and it collapses toward zero. Wealth is measured on the beginning-of-period balance sheet, so liquid wealth and housing refer to the same date, and the bequest is measured at death after the period's saving.

2.3 Target Moments and Fit

The four family-size and timing moments are measured for the 1979–1984 birth cohorts: completed fertility and childlessness for women aged 40–44 in the June 2024 CPS fertility supplement, and the two first-birth timing moments from the NCHS natality microdata for mothers of the same cohorts. These are the youngest cohorts with essentially complete fertility. The environment (prices, incomes, entrant wealth, housing menus) is measured on 2019–2023 cross sections; holding preferences fixed across these adjacent cohorts is a maintained assumption, and a robustness package re-targeting current period fertility rates is planned.

Table 3: Target moments and model fit (provisional and illustrative)

Moment	Source	Data	Model
<i>Family size and timing</i>			
Completed fertility	CPS June 2024, women 40–44	1.918	2.079
Childlessness rate	CPS June 2024, women 40–44	0.188	0.122
Mean age at first birth	NCHS natality, 1979–84 cohorts	25.31	26.40
Share of first births at ages 30+	NCHS natality, 1979–84 cohorts	0.270	0.378
<i>Family housing and tenure</i>			
Housing increment with the first child, rooms	PSID event study	0.664	0.647
Rooms gap, 3+ versus 1–2 children, ages 30–55	ACS	0.368	0.348
Family ownership gap, ages 30–55	ACS	0.168	0.154
Ownership rate, ages 30–55	ACS	0.575	0.410
Mean occupied rooms, ages 18–85	ACS	5.780	6.214
<i>Wealth and bequests</i>			
Aggregate wealth / annual gross labor earnings	PSID 2005–2019	6.873	4.354
Annual bequest flow / aggregate wealth	De Nardi and Yang (2014)	0.0088	0.0091
p_{90}/p_{50} total wealth, ages 76–84	PSID	3.448	2.177

These estimates are provisional and illustrative. They come from a first pass at the revised target system and are reported to show where the specification currently stands, not as a calibration to be used. Four things are wrong with it.

The utility value of children sits at the top of its search region while completed fertility overshoots. Lowering ψ would improve both the level of fertility and childlessness, so the estimate is at a corner rather than an interior optimum, and the fertility block is being pulled in two directions at once. Second, the entry noise remains large, $\kappa_E = 36.3$, and both timing moments miss; the moments meant to discipline it are not doing so, and with first-birth timing this noisy, births respond weakly to prices. Third, the ownership rate is 0.41 against 0.58 in the data, and it moves with the wealth ratio, 4.35 against 6.87: a lower discount factor produces less accumulation and fewer households clear the down payment. Since the down-payment constraint carries the mechanism, this is not a cosmetic miss. Fourth, late-life wealth dispersion is 2.18 against 3.45 with θ_1 near zero, so the bequest motive is not generating the upper tail.

The weights are the first thing to revisit. They mix bootstrap standard errors, working values for the timing rows, and a fixed fraction of the target elsewhere, and that mixture decides which rows win when the fertility block is torn.

Figure 1 compares model homeownership, the share of households with dependent children at home, and mean rooms with their ACS age profiles at the current estimate.




Figure 1: Lifecycle ownership, children at home, and housing consumption at the current estimate. Panel (a) compares model homeownership with the ACS household-head profile. Panel (b) compares the model share of households with dependent children at home with the ACS share of household heads with an own child under age 18, and compares mean rooms in the model and the ACS. The shaded region marks ages 18–45, the fertile window.

References

- Borella, M., M. De Nardi, and F. Yang. 2023. "Are Marriage-Related Taxes and Social Security Benefits Holding Back Female Labour Supply?" *Review of Economic Studies* 90(1): 102–131.
- Davis, M. A., and S. Van Nieuwerburgh. 2015. "Housing, Finance, and the Macroeconomy." *Handbook of Regional and Urban Economics* 5: 753–811.
- De Nardi, M., and F. Yang. 2014. "Bequests and Heterogeneity in Retirement Wealth." *European Economic Review* 72: 182–196.
- Fernández-Villaverde, J., and D. Krueger. 2007. "Consumption over the Life Cycle: Facts from Consumer Expenditure Survey Data." *Review of Economics and Statistics* 89(3): 552–565.
- Flodén, M., and J. Lindé. 2001. "Idiosyncratic Risk in the United States and Sweden: Is There a Role for Government Insurance?" *Review of Economic Dynamics* 4(2): 406–437.
- Greaney, B., A. Parkhomenko, and S. Van Nieuwerburgh. 2025. "Dynamic Urban Economics." Working paper.
- Heathcote, J., K. Storesletten, and G. L. Violante. 2017. "Optimal Tax Progressivity: An Analytical Framework." *Quarterly Journal of Economics* 132(4): 1693–1754.
- Leridon, H. 2004. "Can Assisted Reproduction Technology Compensate for the Natural Decline in Fertility with Age? A Model Assessment." *Human Reproduction* 19(7): 1548–1553.
- Saiz, A. 2010. "The Geographic Determinants of Housing Supply." *Quarterly Journal of Economics* 125(3): 1253–1296.

- Sommer, K. 2016. “Fertility Choice in a Life Cycle Model with Idiosyncratic Uninsurable Earnings Risk.” *Journal of Monetary Economics* 83: 27–38.
- Scholz, J. K., A. Seshadri, and S. Khitatrakun. 2006. “Are Americans Saving ‘Optimally’ for Retirement?” *Journal of Political Economy* 114(4): 607–643.