

Structural review of the housing–fertility OLG model: consolidated reading of two referee reports

Table 5.1 (repository-reading) and ChatGPT Pro, same master prompt

September 15, 2026

Contents

How to read this document	2
Where the two reviews stand	2
Preferences	3
P1. Utility form and the linear child term	3
P2. Sign and role of the cross-partial	4
P3. The first-birth cost ξ	4
P4. Children ever born and children at home	5
P5. Bequests	5
Fertility and children	6
F1. Timing within the period	6
F2. The two extreme-value shocks	6
F3. Stochastic maturation, deterministic adult ageing	6
F4. Dependents of a dying household	7
F5. Exogenous fecundity, chosen attempts	7
Housing and finance	8
H1. The constraint set	8
H2. The underwater rollover	8
H3. The rental cap	9
H4. Supply and rent	9
H5. The landlord	10
H6. The owner premium χ	10
Demography, government, closure	10
D1. Household, person and dependent accounting	10
D2. Population closure	11
D3. PAYGO pensions and the rebate	11
D4. Geography	12

Earnings	12
E1. Permanent types and the persistent AR(1)	12
E2. Earnings risk and fertility	12
Cross-cutting	13
X1. Interactions	13
X2. The smallest defensible set	13
X3. Not on the list	13
Where the mechanism stands (September 17)	14
Follow-up log	15
Combined table	44
Appendix. Verification status of claims made by one reviewer only	46

How to read this document

Two reviewers answered docs/prompts/MASTER_PROMPT_model_review_20260915.md on September 15, 2026. Fable read the repository and verified every code claim against file and line (docs/model/structural_model_review_fable_20260915.md, commit 3f6b5f0b). ChatGPT Pro worked from the prompt, the slides and the ledger text (docs/model/structural_model_review_chatgpt_20260915.md, verbatim copy). For each of the 25 labels this file gives: what the two agree on, condensed; where they disagree, both positions; what only one of them raised, attributed; and the two recommendations side by side. Neither review is edited here; where a reviewer's claim was checkable against code or a local PDF, the check is recorded. Notation follows the prompt: n children ever born, m children at home, $\bar{h}(m)$ the parents' space floor, $e(m)$ the equivalence scale, κ_H the housing taste scale, κ_1, κ_C the fertility taste scales, ξ the first-birth cost.

Where the two reviews stand

Both retain the mechanism and the household block. Both say the core (down payment at purchase, owner-only large units, children raising the value of space, priced fertility timing) is coherent and should not be rebuilt, and both locate the real faults in accounting, closure, supply and finance rather than in preferences. Both say the same three things loudest: the child maturation process is memoryless and must change; the housing stock must have an inherited-stock transition; the population, person and dependent accounts must be written as one contract before any recalibration.

Both independently found the same error in the prompt. The prompt's description of the population closure (imposed historical masses, headship law after 2023) describes an older branch, not the retained one, which propagates households endogenously with the births-over-2.1 queue.

The split is about how much structure to add. ChatGPT is conservative: fix descriptions, normalizations and accounts first; add economic structure (amortization, concave child utility, a child time cost, a rental wedge) only when a specific behavioral failure demands it. Fable is prescriptive: where the reference papers or the data already contradict an implementation choice, change it now (stock-flow supply, an amortizing origination-only mortgage, a size-dependent rental wedge in place of the cap and the owner premium, one

coherent earnings decomposition, a measured child earnings penalty in place of ξ). The disagreements below are all instances of this split.

Sequencing differs. ChatGPT: units and demographic closure first, then preference and credit interpretations, then earnings and geography, then identification. Fable: the two parameter-level repairs first (maturation, earnings coherence), then the household block (mortgage, tenure shock) in one refit, then supply with the terminal condition, then geography with the final refit.

Raised by only one reviewer. ChatGPT: welfare with endogenous fertility needs an explicitly identified population; the down-payment test excludes four years of current earnings; the at-most-one-birth-per-period restriction; the level normalization of a negative bequest utility under child scaling; net depreciation versus paid maintenance; do not reapply the ψ_0 root in counterfactuals; the scale in Scholz–Seshadri–Khitatrakun multiplies utility (verified, and it changes the cost of children from falling to rising in m). Fable: no time cost of children, hence no negative income–fertility gradient; estates have no receiver; no widowhood or downsizing shock; two supply elasticities coexist in code; live depreciation is 1.1 percent, not 2; the tenure-shock value flips the policy’s ownership sign; h_P is at its upper bound; one price per room at every size; the household-formation margin is frozen.

Preferences

P1. Utility form and the linear child term

Agreed. The building blocks (CRRA over goods and space net of a parents’ floor, an equivalence scale, a reduced-form benefit of children at home) are defensible, and the paper must write their rationale. The linear benefit ψm gives no diminishing marginal benefit, and at $\sigma = 2$ with scale exponent 0.7 the flow utility of an additional child at fixed goods and space is convex among households already parenting: the scale burden rises less than proportionally (Fable’s increments 0.234, 0.216, 0.203 for the first three children). Both stress that this does not prove corner completed fertility; it places the stopping margin on dynamics, constraints and the taste shocks rather than on preferences.

Disagreement: what to do. *ChatGPT:* keep provisionally. Describe the object as parental utility over equivalized living standards plus the benefit of children at home, and do not introduce concave child benefits or a different cardinal scaling until a specific behavioral failure justifies it. *Fable:* author decision. Position 1, keep linear and state that κ_C , the fecundity decline and the window carry the intensive margin; position 2, a concave $\psi v(m)$ with externally fixed curvature (Sommer 2016 and Barro–Becker put curvature in the child term), which makes the intensive margin an economic object. Discriminating observable: number of children among mothers by household income and the three-plus share.

Only ChatGPT, verified. Scholz, Seshadri and Khitatrakun’s objective is $E \sum_j \beta^{j-S} n_j U(c_j/n_j)$ with $n_j = (A_j + 0.7K_j)^{0.7}$ (their pp. 615 and 619, checked in the local PDF): the scale multiplies utility as well as deflating consumption. The model uses $U(c/e)$ only, so it imports the scale’s shape, not their objective, and with endogenous family size the difference matters for behavior and welfare. Do not multiply by e merely to match the citation; that would be a substantive change. Becker–Lewis does not derive a linear ψm .

Consolidator’s note on what the verified form implies. Under the SSK objective at $\sigma = 2$ the flow utility is $-e(m)^2/X$, so the cost of successive children is 0.522, 0.580, 0.630 in units of $1/X$, rising in m , whereas under the model’s $U(c/e)$ it is 0.234, 0.216, 0.203, falling. Adopting the cited paper’s actual objective therefore turns the convexity that both reviewers flag into a diminishing net benefit of children without any

change to the child term, at the price of a larger first-child cost and a recalibration of $\psi, \xi, \kappa_1, \kappa_C$. This is a third option for the author decision below, not a recommendation from either review.

Only Fable. Doepke and Kindermann (2019) are the linear precedent: a one-time linear utility from each birth, paired with a fixed goods cost, a fixed utility cost and a wage-scaled time cost per child, so their marginal cost of children is not concave. The scale exponents 0.7 and 0.7 are hard-coded literals (`solver.py:2278`); ψ multiplies children at home (`solver.py:2254`); $h_P = 2.3$ sits at its upper bound with the per-child rooms slope restricted to zero.

Recommendations. Fable: Decide (low; blocks). ChatGPT: Keep (low; blocks, the cardinal interpretation must be chosen).

Follow-up, September 16. Reading done: `child_cost_utility_lit_review_20260916.md` (PDF under `output/pdf/`), forms read from 13 PDFs. Every fertility-choice paper on disk has a marginal cost of children that is constant or rising in the number of children; ours is the only one where it falls, because the deflating scale with exponent 0.7 under $\sigma = 2$ is used without curvature elsewhere. Linear benefits exist (Doepke–Kindermann) but always with linear or wage-proportional costs. The income gradient comes everywhere from a wage-proportional cost, which we lack. Three candidates written out: S1 concave $v(m)$; S2 the SSK weighting $e U(X/e)$; S3 a wage-proportional child cost plus a per-child floor. Lean S3, S2 as a one-line diagnostic. **Status: author decision pending; to be tested in the sandbox.**

P2. Sign and role of the cross-partial

Agreed, fully, on the mathematics. The prompt’s worry is reversed: with $u = \mathcal{C}^{1-\sigma}/(1-\sigma)$, $u_{sm} = \mathcal{C}^{-\sigma}(\mathcal{C}_{sm} - \sigma \mathcal{C}_s \mathcal{C}_m / \mathcal{C})$, so with $\mathcal{C}_m < 0$ the CRRA curvature strengthens a positive cross-partial (ChatGPT’s derivation). For the implemented aggregator with a parenthood-only floor, $\mathcal{C}_{sm} < 0$ for $m \geq 1$ because the scale lowers $\mathcal{C}_s$, while $u_{sm} = (\sigma - 1)(e'/e)u_s > 0$ at $\sigma = 2$ (both). The slide states the restriction on $\mathcal{C}$; it must be restated on u , with finite differences because children are discrete (both; ChatGPT cites `latex/september_14_presentation.tex:127–145, 414–428`). The marginal rate of substitution $u_s/u_c = \frac{1-\alpha}{\alpha} \frac{c}{s-h(m)}$ does not depend on $e(m)$ (both): additional children do not shift the intratemporal space–goods trade-off; only the first child does, through the floor; larger families occupy larger homes only through income, saving and selection.

Only Fable. Both channels raise u_s only when $\sigma \geq 1$; at $\sigma < 1$ the scale channel reverses. At fixed expenditure the floor moves rooms by $\alpha \Delta \bar{h}$: 1.69 rooms for an unconstrained renter against 1.01 realized, the gap being the constraints (cap, sale cost). The scale multiplies future marginal utility by 1.23 for a first child, which is where anticipatory saving for a down payment comes from.

Only ChatGPT. A positive u_{sm} does not establish the sign of the optimized fertility response to prices; three properties must be kept apart (marginal utility of space, relative housing demand, fertility response to cost). Add a per-child space component only if the paper intends and supports a separate requirement per additional child.

Recommendations. Fable: Keep (none; no). ChatGPT: Change (low; no). Same substance: restate the property on u .

P3. The first-birth cost ξ

Agreed. ξ shifts the first-birth index by $\pi_a \xi$; κ_1 rescales that index and its response to costs; the two are not interchangeable, and ξ is a childlessness lever by construction, which is fair rather than an objection. The

identification mapping on the slide is a reading of the parameter table, not a verified Jacobian (M25).

Disagreement: keep or replace. *ChatGPT:* keep, as an explicit fixed utility cost of entering parenthood, not an unspecified goods expenditure. Run the zero- ξ restriction holding the ψ_0 root; removal may preserve the childlessness mean but there is no guarantee it preserves the joint age-wealth pattern of first births. *Fable:* author decision, lean change. Replace it with a measured per-period earnings cost of children at home (a child penalty on y^d while $m > 0$, Kleven et al.), which scales with the wage and yields the negative income gradient of fertility; keep ξ only if childlessness still needs it. The nearest precedent, Doepke and Kindermann's fixed utility cost per child, sits beside a time cost and cannot be borrowed alone.

Observables. Both: first-birth hazards by age and liquid wealth together with completed childlessness. Fable adds childlessness by permanent income or education, on which ξ predicts a flat gradient and an earnings penalty a negative one.

Recommendations. Fable: Decide, lean change (low; blocks). ChatGPT: Keep (low; no).

P4. Children ever born and children at home

Agreed. Keep both. A household with $m = 0$ may be childless or an empty nest, and the two face different first-birth costs, scales and remaining opportunities; the three ever-born moments need n . Fable quantifies why the confusion is not rare: under memoryless maturation 22 percent of first children leave within four years.

Implementation notes. ChatGPT: store only feasible combinations $0 \leq m \leq n$ and share continuation computations outside fertile ages. Fable: drop n from the state after the last fertile age, since nothing after 46 depends on it once the bequest is child-blind.

Recommendations. Keep (both; low or none; no).

P5. Bequests

Agreed. The child-blind warm glow is the right object (De Nardi; Kopczuk and Lupton); do not restore scaling in the number of children; the estate must be net bequeathable wealth at a consistent death-date price, net of mortgage and any liquidation cost. Fable verified that the code values the house gross of the 6 percent selling cost (`solver.py:2455–2459` against `solver.py:2442`), so dying in the house is cheaper than selling it.

Only ChatGPT. For $\sigma > 1$ the unnormalized B is negative, so multiplying it by the number of children injects a fertility incentive through the level normalization; child-specific altruism would need recipients and an estate-division rule. De Nardi's framework does not imply a preference for bequeathing housing over financial assets; old-house retention must come from actual frictions.

Only Fable. The estate has no receiver: it is a utility argument and the wealth leaves the economy, a choice Greaney et al. also make for their entrants. For a paper about intergenerational allocation, route estates to households aged 45–65 as a lump sum; this also relieves the wealth-to-earnings undershoot that pushes β to its cap. The larger omission for the policy is the absence of every downsizing force (widowhood, care, health; Venti and Wise 2004, to verify): the old release large homes only at death, so a holding tax moves them only along the intensive margin at a 6 percent cost. Author decision on an age-rising forced-move shock.

Recommendations. Fable: Change the two implementation points (medium; blocks). ChatGPT: Keep the net-estate motive (low; blocks, the net-estate accounting must be settled). Substance close; the receiver

question is Fable’s alone.

Fertility and children

F1. Timing within the period

Agreed. Keep. The sequence (fertility shock, attempt, conception, housing shock, housing choice) is an information structure, not a solution device: households adjust housing within the four years in which a birth occurs, which is the object the PSID rooms response measures. Committing to a house before conception is a different problem and belongs in a targeted robustness run, not in the baseline. Do not read the sequence as households waiting until delivery to move.

Only Fable. The September 6 nested-choice experiment found the joint-versus-sequential difference at fixed continuation values to be of order 10^{-6} births per household; it did not test commitment.

Only ChatGPT. Observables should include moves before unsuccessful or postponed fertility plans where they can be seen.

Recommendations. Keep (both; none or low; no).

F2. The two extreme-value shocks

Agreed. The fertility shocks are economic primitives (unobserved motivations, identified by the level and age spread of first births). The housing shock at $\kappa_H = 0.005$ is a search-bound artifact with no dedicated moment, not an estimate, and it should leave the estimated vector. The extreme-value location normalization must be stated: Fable notes the code’s $\kappa \log \sum \exp$ is exact only for mean-zero shocks; ChatGPT notes that omitted scale-dependent constants are not harmless when scales differ by reproductive history.

Disagreement: what replaces the estimated κ_H . *ChatGPT:* fix it as a numerical approximation with a stated limiting target as smoothing vanishes; check menu dependence, since K equal-valued alternatives add $\kappa_H \log K$ to the log-sum and grid refinement must not create a taste benefit. *Fable:* either $\kappa_H = 0$ with the kink handled at the distribution level, or estimate it against a tenure-transition moment (the PSID four-year Brier score dropped on July 24). It cannot stay a free numerical choice: the August 20 vintage at $\kappa_H = 0$ gives a positive ownership response to the tax and the September 1 vintage at 0.010 a negative one. Greaney et al. choose tenure and size deterministically, so a product-level taste shock is not “DUE-style”.

Recommendations. Change (both; Fable low to medium, ChatGPT low; blocks).

F3. Stochastic maturation, deterministic adult ageing

Agreed. The asymmetry is fine; memorylessness is the fault. With per-period exit μ , duration has no minimum and no maximum: newborns enter the first draw and dependents persist into parental old age (Fable: $\mu = 4/18$, 22 percent of newborns gone within four years, 8 percent of children born at parental age 25 still home at 65). Independence across siblings is secondary. If the process stays, ChatGPT would rename it “departure from dependency” and stop calling exits maturation.

Disagreement: the fix. *ChatGPT:* a bounded stochastic dependency duration through a short child-age-stage or birth-history state. Stochasticity can remain, memorylessness need not. This reopens the author’s recorded decision to set bounded durations aside and is, in ChatGPT’s words, the largest substantive household-state decision in the package. *Fable:* no new state. Let the exit hazard depend on the parent’s age, $\mu(a)$, low through the fertile window and rising to one by 62; since the last birth is at 46 and mortality starts at 66,

every child leaves before any parent can die. Add a one-bit newborn flag that exempts a newborn from the first draw. An Erlang-2 law with the same mean cuts newborn exits to 7 percent but leaves 25 percent at home after 24 years, so a two-stage clock alone does not fix the tail. *Trade-off between the two*: the parent-age hazard zeroes the orphan flow and thins the tail at zero cost in state space, but it mis-times departures for children born early in the window, who stay until the parent’s hazard rises; the child-stage state gets timing right at the cost of state space and the reopened decision.

Observable (both). Dependents at home by parent age and time since birth (ACS), which the model overstates above 60 and understates at 25–35.

Recommendations. Change (both; Fable low to medium, ChatGPT medium; blocks).

Follow-up, September 16: children at home by parent age, model versus ACS 2005–06 (sandbox baseline, fixed ψ ; figure and table in output/model/sandbox/dependents_by_parent_age/). The share of households with any child at home matches the ACS at every age. The mean number does not: the model peaks at 0.89 children at ages 34–38 against 1.44 in the ACS (1.39 counting only children under 18), and is 0.75 against 1.18 at 30 and 0.85 against 1.29 at 42. Above 58 the model’s dependents (0.31 at 58, 0.19 at 66) track the ACS count of own children of any age (0.30, 0.17) but not the count of minors (0.06, 0.01): the memoryless tail is reproducing adult children who live with their parents, not dependents. Reading: the constant hazard empties the nest too early and too slowly at once. The young-age shortfall of roughly 40 percent is larger than the sandbox’s completed fertility gap (11 percent) and sits on the space-demand margin the paper is about. The doubt that memorylessness is harmless is answered: it is not. It also says which fix: the young end needs a minimum duration (newborn flag or a child-stage state); the old end needs an age-dependent exit. **Status: author decision between the parent-age hazard plus newborn flag and a child-stage state.**

F4. Dependents of a dying household

Agreed. If F3 removes dependency before any parental mortality, the problem disappears and no care pool is needed; that invariant should be established rather than a repair added. Joint death changes the mortality model rather than the bookkeeping; reassignment adds state and fertility effects.

Disagreement, conditional. If the current dependency process stays, ChatGPT prefers an aggregate care pool (goods and housing needs, funding, an exit rule, no second entry stream) over joint death or reassignment; Fable treats all three as symptom repairs and would build none.

Only ChatGPT. The 1.2 percent flow was measured on an earlier checkpoint, not the retained baseline; and because the recorded-birth queue already preserves those children’s future entry, the established gap is their care and housing, not a loss of persons.

Only Fable. Verified code: `run_e5f_open_population_transition.py:847–851, 887`; the 2023 measurement is 1.01 percent of dependents.

Recommendations. Change (both), through F3.

F5. Exogenous fecundity, chosen attempts

Agreed. Keep the split (Sommer 2016’s absorbing infertility shock; Doepke–Kindermann’s birth probability with natural fecundity and imperfect birth control; de la Croix–Pommeret’s fitted schedule). π_a must be success conditional on the modeled attempt, not an observed birth rate that already contains choices.

Only ChatGPT. The restriction of at most one birth per four-year period imposes spacing and limits catch-up after delayed first births; losses in completed fertility must not all be attributed to biology. Attempts are costless, so estimated attempt probabilities at very low π_a are not literal. Define the four-year attempt and map biological evidence to that horizon.

Only Fable. The live schedule (0.02, 0.134) gives 0.97, 0.90, 0.83, 0.71, 0.50 at ages 22, 30, 34, 38, 42, within three points of Léridon’s four-year figures, but was chosen by eye. Every birth is chosen, whereas about 45 percent of U.S. pregnancies were unintended in 2011; an exogenous birth hazard for non-attempters (author decision) would lower the price elasticity of total births by roughly the intended share.

Recommendations. Keep (both). Fable: low, does not block. ChatGPT: medium, blocks (horizon and birth-count mapping).

Housing and finance

H1. The constraint set

Agreed. The current formulation lets an owner re-borrow up to the collateral limit each period against an unchanged house, so “no refinancing” is not literally true (Fable: `kernels.py:979–985`, at no cost). The slide’s down-payment inequality is wrong for stayers: its right-hand side is zero, leaving a net-equity restriction, whereas the code tests only on trades (ChatGPT’s reading; Fable: `kernels.py:640–662`). The gross-return factors are not an error; housing trades precede current earnings by design. ϕ stays.

Disagreement: how much to add. *ChatGPT:* change the formal description and the purchase-condition logic only. State the down payment as a restriction on purchase transactions; specify incumbent debt and new borrowing separately; describe the signed-asset formulation as a reduced-form collateral technology. Do not add rate risk, amortization and HELOCs at once; preserve the parsimonious core. *Fable:* add a required amortization share, apply the collateral test at origination only, forbid cash-out for stayers, and add a payment-to-income test if the Coven comparison is to be like-for-like. Free cash-out mutes Coven’s cash-flow channel (the old pay the holding tax out of equity) and removes the buffer-stock motive of young owners. Verified: Greaney et al. have the model’s current block (origination LTV, no amortization); Kaplan, Mitman and Violante and Coven et al. both have amortization and payment-to-income limits.

Only ChatGPT. The down-payment test excludes four years of current earnings from purchase liquidity. At a four-year period this is a strong assumption and, in ChatGPT’s view, the more immediate concern.

Recommendations. Change (both; medium; blocks), with different content.

H2. The underwater rollover

Agreed. Grandfathering existing debt after a price fall is right; Kaplan, Mitman and Violante distinguish origination limits from limits on existing mortgages. The age taper is ad hoc, must be exposed in the paper, and must not be deleted from the code just to match the slides.

ChatGPT. Replace or rationalize the taper with a contract rule that distinguishes scheduled outstanding debt from new borrowing; inherited negative equity is not permission to originate unsecured credit; specify the creditor and settlement rule for sale shortfalls and terminal unpaid debt.

Fable. The taper disappears under H1’s amortization; otherwise disclose it as an assumption. Code notes: one taper array doubles as the unsecured-credit scale (inert at $\lambda_d = 0$); sale shortfalls carry into the renter state and decay under the same taper (`solver.py:3471–3474`).

Recommendations. Change (both; Fable low, ChatGPT medium; blocks).

H3. The rental cap

Agreed. The cap is an approximation of segmentation, not an institutional fact. Greaney et al. cap rental size and calibrate the level from the size distribution (Fable verified: 5.39 against 10.78, the 90th percentile of rentals over the 10th percentile of owner units), so “6 rooms, Greaney et al.” should become that recipe measured in the AHS. The paper needs a soft-tail alternative (size-specific rental supply or a steep size-dependent premium) and the joint size–tenure distribution to discipline it; the question is whether the result needs literally zero access or survives expensive, limited access.

Disagreement: the baseline. *ChatGPT:* a hard-cap baseline can remain, with the soft-tail comparison as robustness. *Fable:* make the size-dependent wedge the baseline, $r_t(h^R) = r_t + w(h^R)$, calibrated to AHS renter shares by size; the wedge is Kaplan, Mitman and Violante’s landlord operating cost (verified) and Henderson–Ioannides’s rental externality (from memory). Evidence for the wedge over the wall: 7.5 percent of renters live in seven-plus-room units and renters hold 8 percent of that stock; 31.8 percent of renters aged 25–45 sit exactly at the cap, a mass point that creates the kinks κ_H smooths.

Recommendations. Change (both; medium; blocks).

H4. Supply and rent

Agreed. Replace the instantly reversible schedule $H_0 P^\eta$ with an inherited stock and a construction margin (ChatGPT: $I_t \geq 0$; Fable: $I_t = I_0 P_t^\eta$, the reference paper’s own block, with Kaplan, Mitman and Violante’s construction sector as a second precedent; Coven et al. use the same static form as the model, so the Coven comparison is like-for-like either way). Keep the forward-looking user cost; it is coherent under competitive landlords and a deterministic continuation. Both reviews said the 0.63 elasticity has no derivation from Baum-Snow and Han; the local PDF partly corrects both. Their headline averages are 0.5 for floor space and 0.3 for units across urban neighborhoods (abstract), 0.42 and 0.35 at tract level (pp. 1898–1899), and a metro-aggregate floor-space elasticity of 0.61 to 0.63 appears under their linear-IV specification (p. 1937). So 0.63 is a metro-level floor-space number from the paper, not a neighborhood or housing-unit elasticity; the deck should cite that page and say which object it is.

Only ChatGPT. Define net depreciation consistently with the maintenance owners already pay, so the same physical depreciation is not counted twice. Under successive surprises, use the forecast known at each date, not prices generated by later news. If H^s means occupied services rather than stock, account for the inherited stock and vacancies.

Only Fable. Impact arithmetic: the tax’s 17 percent price fall removes about 11 percent of housing within one period under the static rule, while depreciation allows at most 8 percent in four years; this is the “services fall” the September 1 note misread. Long run: with a level supply curve and a population going to zero, $P \rightarrow 0$ and no per-household stationary limit exists, which is why there is no terminal steady state (ties to D2). Two elasticities coexist in code, 1.75 in the profile chain and 0.63 as a transition override (e1_profile.py:39; the transition manifest). Live depreciation is 1.1 percent a year (run_e1_chain.py:389), not the slide’s 2 percent.

Recommendations. Change (both; medium). Fable: blocks the transition and H_0 only. ChatGPT: blocks.

H5. The landlord

Agreed. The closure is an outside competitive rental company earning the exogenous return, Greaney et al.’s REIT with the same rent rule (Fable verified). Its stock ownership, financing, tax payments and cash flows must be written; resident welfare excludes nonresident owners unless counted; rent payments and unexpected capital losses do not vanish because expected excess returns are zero. Do not route them through the property-tax rebate (ChatGPT); report them as accounting lines (Fable). Fable verified that the tax base includes the rented stock and that no landlord code exists.

Only ChatGPT. Domestic household landlords (Sommer–Sullivan–Verbrugge) are an admissible but non-equivalent alternative closure.

Only Fable. In the June 2025 Coven draft on disk, rent is a user-cost no-arbitrage price with the tax inside, so the tax passes to renters there as here; the September 1 note attributes a different rental block to the August 2026 version, which is not on disk, so the difference is unverified either way.

Recommendations. Fable: Keep and state (low; no). ChatGPT: Change the documentation and accounts (medium; blocks). Same substance.

H6. The owner premium χ

Agreed. The code applies the premium as $\chi(h - \bar{h}(m))$ (`kernel.py:957–970`) and the slide as $\chi h - \bar{h}(m)$; these differ, the latter lowering the owner’s physical minimum to $\bar{h}(m)/\chi$. Decide which is meant and make code and exposition agree.

Disagreement: keep or replace. *ChatGPT:* keep χ as a legitimate reduced form; do not call it a missing financial motive, since ownership already provides an asset position and collateral services; do not use it as an unrestricted ownership-fit residual. Verified in the local PDF: Greaney et al. define housing services as $\chi h + h^r$, call χ the non-pecuniary benefit of ownership, and calibrate $\chi = 1.0506$ to a homeownership rate of 0.542 (their pp. 10–11 and Table 1); with $\chi = 0$ no one owns (p. 42). The model’s $\chi = 1.048$ is the same device at the same value, which supports ChatGPT’s position; note that the reference form is χh , the slide’s, not the code’s $\chi(h - \bar{h})$. *Fable:* in the model renting and owning the same size cost the same user cost, so χ is the only reason to own and stands in for the tax advantages of owning, the rental externality and the rent-risk hedge; fold it into the intercept of H3’s rental wedge, since the ownership rate identifies only the sum of a premium and a wedge, and the price–rent ratio separates them.

Recommendations. Fable: Decide, lean change (medium; blocks). ChatGPT: Keep (low; blocks, the physical-space interpretation).

Demography, government, closure

D1. Household, person and dependent accounting

Agreed. Write one accounting: births to queue to entrants, survivors to maturation, deaths to estates, units stated, estates reconciled exactly once. The person, household and dependent counts are not linked in the current implementation; person and head aggregates pass their separate checks while the dependent-child link is unresolved. This is a prerequisite for population and intergenerational-allocation claims.

Disagreement: what the household is. *ChatGPT:* a unitary household need not contain two adults, and the two-adult base of the scale does not establish a two-person unit. Use an unnormalized household measure with explicit person-stock identities: $D_{t+1} = D_t + B_t - M_t - \Delta_t^D$ for dependents and $A_{t+1} = A_t + M_t - \Delta_t^A$

for adults, with household formation a separate mapping from adults to households, not one household per maturing child. *Fable*: three conventions already make it an implicit couple (head-plus-spouse earnings, the two-adult scale base, entrants equal to births over 2.1), so say so; the missing margin is then dissolution by widowhood, which is how most old households become one-person households in large homes (author decision).

Only Fable. Entrants arrive twenty years after birth (four waiting slots), as childless renters with PSID 18–24 wealth ratios; the 5 percent leakage inside 2.1 is an external constant while the model’s own orphan leakage is another 5 percent of births; the 2007 stationary normalization forces a replacement population while the childlessness and one-child moments are for the 1960–66 cohort and the timing moments are 2003–06 period statistics.

Only ChatGPT. Physical market clearing must use $h + h^R$, not $\chi h + h^R$ (the code does clear in physical units, per *Fable*’s verifier); orphan transfers are internal movements, not person losses.

Recommendations. *Fable*: Keep and write (low; no). *ChatGPT*: Change (medium; blocks).

D2. Population closure

Agreed. The prompt describes the wrong branch: the retained branch propagates households endogenously with the births-over-2.1 queue, no imposed masses, no headship law, no migration (both, *ChatGPT* from M26, *Fable* from spec.json). The counterfactual must keep the inherited 2023 distribution and pre-2023 birth pipeline and let policy births affect entry only at the later dates. Matching completed fertility to 2.1 in 2007 does not itself establish replacement under the model’s own survival and formation rules.

Complementary, not conflicting. *ChatGPT*: choose one demographic interpretation, a literal person model in which survival, entry and formation agree, or an abstract reproductive household with a conversion normalization, and never present converted household counts as persons. *Fable*: the binding problem is the terminal condition. With zero migration and completed fertility below 2.1 there is no positive stationary population ($B/E \leq 0.87$ at any price), and with level supply no per-household stationary limit either, so the transition solver’s terminal state is a fiction. Two coherent options: a policy-invariant outside inflow (the August 6 closure), or the author’s zero migration with a balanced-contraction terminal condition, which requires H4’s stock-flow supply. The headship law, where used, freezes household formation at 2023 rates.

Recommendations. *Fable*: Decide (low or high). *ChatGPT*: Change (high). Both: blocks the transition and policy, not the initial calibration.

D3. PAYGO pensions and the rebate

Agreed. Keep. Fixed payroll tax with the benefit adjusting each period is coherent; the equal rebate is Coven’s rule (*Fable* verified) and defines a particular redistribution experiment rather than an accounting detail. Write the base and recipients: $T_t \mathcal{H}_t = \tau_t^p P_t (H_t^O + H_t^R)$ when both stocks are taxed at the rate in the user cost (*Fable* verified that they are). Report capitalization, net tax payments, rebates and pension adjustment separately rather than inferring incidence from age.

Only ChatGPT. Any care-pool funding taken from the receipts changes the full-rebate experiment and must be identified.

Only Fable. The alternative closure (fixed replacement rate, adjusting tax) would load ageing onto the young; the rebate enters the period budget after the housing choice, not the down-payment test; in a shrinking population pensions fall and that is a downsizing force the model otherwise lacks.

Recommendations. Keep (both; none or low).

D4. Geography

Agreed. A national pooled benchmark: remeasure the housing moments nationally under the same sample rules and reconcile prices, supply, earnings and demographic units with that population. The metro alternative needs metro-scoped fertility and demographic evidence and raises migration and sample-support questions. Do not relabel the existing estimates.

Recommendations. Change (both; medium; blocks).

Earnings

E1. Permanent types and the persistent AR(1)

Agreed. Reconcile the income definition, population, taxes, age effects, variances and autocovariances before deleting states; a persistent-only alternative must be refitted to the same evidence; the four-year conversion uses ρ^4 with accumulated innovations. Permanent heterogeneity is not contrary to the mechanism.

Difference of evidence, not of principle. *ChatGPT:* the ledger does not establish compatibility or duplication; keep provisionally and remove only if a comparably disciplined persistent-only process reproduces the evidence. *Fable, from code:* the live AR(1) is Floden–Lindé scaled to after-tax units (annual $\rho = 0.9136$, innovation variance 0.0426 times $(1 - 0.181)^2$, stationary log-variance 0.173) while the fixed effect is the E6b PSID gross-earnings variance 0.393; persistent-plus-fixed is 0.566 in the model against 0.725 in the decomposition the fixed effect came from, in different tax units (0.264 after the same scaling). Not double counting; an incoherent sum. Fix: use the E6b decomposition for all three components with one tax treatment, then run the removal comparison. The types were added on July 27 to reach the old-age wealth tail; the age profile e_a has no source.

Recommendations. Fable: Change (low; blocks). ChatGPT: Keep (medium; blocks, reconcile the evidence). Same procedure; Fable supplies the finding.

Follow-up, September 17. “Yeah I think we probably have to get rid of that...” The author tentatively leans toward removing permanent earnings types and considering persistent-plus-iid-transitory earnings risk, as in Boar–Gorea–Midrigan and Sommer; final decision pending author reflection. This is a note only, not an adopted specification or recalibration mandate: any comparison must preserve empirical identification and reconcile the earnings evidence, tax units, wealth gradients and fertility gradients.

E2. Earnings risk and fertility

Agreed. Keep. The channel already exists through continuation values and the option value of waiting; no second uncertainty term. Do not assert that every mean-preserving spread lowers fertility; the object is the difference between having a child and waiting. Fable adds that at four years the AR(1) persistence is 0.70, which mutes the channel.

Recommendations. Keep (both; none or low; no).

Cross-cutting

X1. Interactions

Shared. F3 with F4 (one maturation change may remove the orphan problem); rental segmentation with credit and χ ; landlord ownership with taxes and welfare; the scale normalization with $\psi, \xi, \kappa_1, \kappa_C$.

Only ChatGPT. Keep the ψ_0 root as an initial-calibration restriction but never reapply it in a counterfactual to restore 2.1; that would delete part of the response being studied. Moment-to-parameter labels are not identification.

Only Fable. H4 with D2: stock-flow supply is what makes a balanced-contraction long run possible. P1, P3 and E1 meet at the income gradient of fertility (E6b's reversed childlessness gradient is their joint symptom). P5 with β at its cap. D4 with every estimated parameter.

Sequencing, both orders. ChatGPT: units and demographic closure; then preference and credit interpretations; then earnings and geographic scope; then identify the revised vector. Fable: F3 and E1 (parameter-level, cheap); then H1 and F2 in one refit; then H4 with D2's terminal condition; then D4 with the final refit.

X2. The smallest defensible set

Shared package. Reconcile dependents, persons and entrants (D1, F3, F4); formalize purchase credit and incumbent debt (H1, H2); specify landlord, estate and fiscal accounts (H5, P5, D3); one empirical population with a compatible earnings process (D4, E1); an inherited housing stock (H4); the shock normalization and the housing-smoothing treatment (F2). Both retain the separate child states, the attempt-and-success structure, the borrowing core and the warm-glow bequest.

Where the packages differ. ChatGPT classes concave child benefits, commitment-before-birth timing and richer mortgage instruments as conditional extensions, and calls bounded dependency the largest household-state decision. Fable puts amortization with origination-only borrowing inside the package, names the time cost of children as the largest omission, and lists the items that need explanation rather than change (P2, P4, F1, F5, H5, D1, D3, E2). Both agree that housing adjustment around births does not validate the reverse effect of housing costs on births.

Author decisions the package leaves open (union of both lists). P1 linear versus concave; P3 keep ξ or replace it with a measured penalty; F3 parent-age hazard versus child-stage state; F4 care pool only if F3 is not adopted; F5 unintended births; H1 parsimonious description versus an amortizing origination-only contract; H3 with H6, cap plus premium versus one wedge; D1 the household unit and widowhood; D2 outside inflow versus balanced contraction; P5 estate receiver and a downsizing shock.

X3. Not on the list

Only ChatGPT. (1) Welfare with endogenous fertility: define the population (for example households alive in 2023 and their remaining utilities), report future cohorts separately, and never sum over different numbers of households without explicit weights and a utility-level convention; more births are not a welfare gain by themselves. (2) Omitted nonhousing costs of children may be absorbed into housing and fertility preferences; an externally disciplined goods and time cost is the targeted alternative. (3) The feasible set must enforce rent-or-own; annual flows and four-year utility and discounting conventions must match; the one-birth-per-period restriction must be acknowledged. (4) A discount factor at its cap is a diagnostic of fit or identification, not proof of invalidity.

Only Fable. (1) No time cost of children: the cost of a child is a homothetic scale plus an absolute floor, so the model cannot produce the negative income–fertility gradient; both Sommer and Doepke–Kindermann carry the time cost. (2) No household-formation margin; the headship law freezes it. (3) No widowhood or health-driven downsizing. (4) The taste shocks must be declared mean-zero type-I extreme value (no Euler constant in the log-sum, `utils.py:194–198`). (5) Depreciation compounds while the property tax is four times annual (`parameters.py:163–165`). (6) $\beta R > 1$ at the cap because the old have no expense risk, no return above R_b and no inheritances. (7) Owner services are $\chi(h - \bar{h})$ in code, $\chi h - \bar{h}$ on the slide. (8) h_P at its upper bound with the per-child slope restricted to zero, so the two rooms moments pull one parameter in opposite directions. (9) One price per room at every size, against hedonic evidence. (10) The fecundity schedule is set in `run_e1_chain.py:380`, invisible from the calibration layer.

Both, with different weight. The time cost of children: Fable’s largest omission, ChatGPT’s conditional alternative.

Where the mechanism stands (September 17)

Written after the first sandbox tests, before any refit. Everything here is a steady state at the retained parameters with the child-preference level held fixed, so it is about direction and size, not about the slides’ levels.

The mechanism on the slide is not what the model does at this calibration. The slide says: children raise the value of space, large homes need a down payment, so house prices shift the timing of births. In the model, removing the down payment raises ownership among 30–55-year-olds by ten points and leaves completed fertility and the age at first birth unchanged. Giving households an unsecured credit line while keeping the down payment leaves ownership unchanged and raises completed fertility by 0.08 with first births a year earlier. The down payment governs ownership; liquidity governs births.

Why, most likely. A renter can have six rooms and the parents’ space floor is 2.3 rooms, so a family with one or two children fits in a rental. The owner-only sizes start at eight rooms and matter for three-plus families. The down payment therefore never blocks the space a first or second child needs; what blocks an early birth is that a renter cannot borrow at all and must save the child’s cost in advance. That is a Sommer-type precautionary channel, and it is real in the model, but it is not the housing-collateral story.

Two ways this can still be a calibration artifact, both under test. (1) The model gives households too much space: mean rooms 6.4 against 5.6 in the data, a first-birth rooms response of 0.98 against 0.72. If space were scarce at the data’s level, the family-size unit might fall into the owner-only range for more households and the down payment might bind on the family margin. This is the author’s prior and it is being tested by lowering the supply scale until mean rooms hit the data and repeating the no-down-payment comparison (task `TASK_muse_constrained_households_20260917.md`). (2) The baseline is the extreme liquidity case: zero unsecured credit for renters, and a down-payment test that ignores four years of current earnings. Real households have some unsecured credit (Kaplan and Violante calibrate a limit near three quarters of a quarter’s income). Part of the fertility response to credit may be the artifact of starting from zero; a modest unsecured line belongs in the baseline decision for H1.

Result of the scarce-space test (September 17, evening). With the supply scale lowered until mean rooms equal the data (5.56, price up 4 percent), the share of family-forming households (ages 26–38, zero or one child at home) whose housing choice changes when the down payment is removed falls from 4.5 to 2.8 percent, and completed fertility is flat with and without the down payment (1.840 against 1.835). Scarcer space lowers fertility a little (1.872 to 1.840, about 0.4 percent per percent of price) through the cost of space, not through the collateral requirement. So the rooms misses are not what hides the mechanism: at the data’s

space level the down payment still governs ownership (0.437 to 0.532) and not births. The author's prior is rejected on this test. Tables: `output/model/sandbox/constrained/README.md`.

What survives either way. The model, the calibration machinery, the transition and the property-tax experiment all stand. Two clean results already exist: easier mortgages raise ownership and not births; easier unsecured credit raises births and not ownership. The property-tax channel to young families runs through ownership and the allocation of large homes, so its fertility effect will be small unless the space margin binds, which is what the scarce-space test decides. If the down payment binds once rooms are right, the priority is the rooms calibration and the slide's sentence is kept. If it does not, the mechanism paragraph is rewritten around liquidity and the rental cap, and the paper's claim changes rather than disappears.

Decisions this touches. H1 (unsecured credit in the baseline), H3 (the cap level is now load-bearing for the mechanism, not only for the ownership fit), P2 (a per-child floor would move the space margin to the second child), and the wording of the mechanism on the slides.

Follow-up log

Every study that acts on a label above gets a dated entry under that label (quoted block) and a line here. This file is the single tracking document for the structural decisions; do not open parallel notes without an entry.

Date	Label	What was done	Pointer	Status
2026-09-16	P1	Literature reading on the benefit and cost of children; three candidate specifications	<code>docs/model/child_cost_utility_lit_review_20260916.md</code>	Author decision pending
2026-09-17 (evening)	H1, H3, X3	Who-is-constrained table and scarce-space test (Muse): down payment binds for 4.5% of family-forming households at baseline, 2.8% once mean rooms equal the data; fertility flat with and without the down payment in both cases; scarcer space lowers fertility through cost (1.872→1.840). The rooms misses do not hide the mechanism	<code>output/model/sandbox/constrained/README.md</code>	Prior rejected; mechanism wording must change

Date	Label	What was done	Pointer	Status
2026-09-19 (09:30)	P5, X2	Torch ψ-root runs collected (jobs 17949060, 2.2 h; 17949061, 8.8 h). Estate transfer to ages 45–65 with net-of-selling-cost bequests: loss 2582 (baseline 2607), ψ 0.181, first-birth age 26.4, old tail 4.07 (4.32), ownership 0.446; a small change. All switches together (penalty, mortgage block, wedge, estate, parent-age maturation, log benefit): loss 962, ψ 0.480, childless 0.223, first-birth age 27.6, first births 30+ 0.32, wealth/earnings 4.4, rooms 5.25, three-plus gap 0.58, ownership 0.426, recent-parent gap 0.27. The switches do not add up: the wedge alone (402) beats all together (962); penalty plus maturation overshoot the timing rows and need ψ two and a half times the baseline	output/model/sandbox/torch_psi_root/, output/model/sandbox/switches_root_psi_root_comparison.md	Rows in the situation report; author decisions pending

Date	Label	What was done	Pointer	Status
2026-09-19 (09:30)	H3, H6	Rental cap at eight rooms on the paper's model (local, fixed ψ 3 min, ψ root 12 min). Fixed ψ : completed fertility 1.876 vs 1.872, childless 0.238 vs 0.239, first-birth age 26.94 vs 26.96, ownership 0.31 vs 0.46, recent-parent gap 0.21 vs 0.46. ψ root: loss 489 (2607), ψ 0.189 (0.190), ownership 0.300, rooms response 1.36. On this model the cap moves ownership, not births; the July 1 audit's "TFR +0.071" was on the earlier package and does not carry over. The tenure memo's "strongest fertility lever" line is withdrawn for the paper's model	output/model/sandbox/cap_eight_psi_fixed/, output/model/sandbox/cap_eight_psi_root/	For the tenure discussion

Date	Label	What was done	Pointer	Status
2026-09-18 (13:30)	H1, X3	How many are constrained, from the literature (Opus subagent, lead-verified on the two load-bearing sources): four objects get called “the constrained share”; only the marginal group (choice changes when the constraint is relaxed) is comparable to the model. Haurin, Hendershott and Wachter (NLSY, ages 20–33): 37 percent constrained times a 10–18 point ownership response gives 3.7–6.7 percent marginal; Kaplan, Mitman and Violante: about 3 percent, “very few households are constrained in this way” because they rent the size they want. The model’s 4.5 percent (a stock object) sits in that range. The two-thirds wealth-constrained figures (Barakova et al.) are an accounting object, not the marginal share. Proposed discipline: SCF share of renters 26–38 with liquid assets below 20 percent of the local median price as a	docs/model/evidence_constrained_share_20260918.md	Author to read; the moment needs SCF microdata

Date	Label	What was done	Pointer	Status
2026-09-18 (13:30)	H3, H6	Tenure segmentation evidence (Opus subagent, lead-verified in part): renter share falls from 85 percent of one-bedroom units to 11 percent of five-bedroom (ACS 2024); scarcity of large rentals shows in quantities and unobserved quality, not in posted rent (Halket, Nesheim and Oswald, London; our own within-cell rent-per-room estimate is zero); KMV use the same partial segmentation and, in the memo’s reading of the published version, find it inessential for their results; DUE sets the rental maximum equal to the owner maximum (verified in the local text), so the July cap note’s “DUE rule” attribution is wrong. The claim that the cap is the model’s strongest fertility lever comes from the July 1 audit on the earlier package and is unmeasured on the paper’s model; a cap-at-eight	docs/model/evidence_tenure_segmentation_20260918.md	Author to read; cap-at-eight run pending

Date	Label	What was done	Pointer	Status
2026-09-18 (12:30)	P5, X2	Torch resubmission (Muse): the batch on scratch lacked the corrected_initial/ folder the launcher expects; recreated with the eight files (md5 matched to the local copy); smoke 17948921 completed in 39 s with the four output files; the two long ψ -root specs resubmitted with 12 h limits, estate transfer 17949060 and all switches 17949061, both running. A ten-minute poll collects them into output/model/sandbox/torch_psi_root/ (one folder per spec) when done	output/model/sandbox/torch_psi_root/jobs.json (resubmission block)	Running; rows in the situation report open until collected

Date	Label	What was done	Pointer	Status
2026-09-18 (11:05)	all	Morning: all ten Torch spec jobs and the smoke show CANCELLED with zero elapsed (never ran; cancelled by a user or admin, not this session). Local results complete for everything except the two estate-loop specs. Situation report finalized: docs/model/situation_report_20260918.md and its PDF	output/model/sandbox/torch_psi_root/jobs.json	Resubmit after the cancellation is understood
2026-09-18 (02:45)	P1	Concave child benefit with ψ re-normalized: loss 2111 (baseline 2607), ψ 0.199, first-birth age 25.90 on the 25.98 target, first births 30+ 0.223, three-plus rooms gap 0.15, ownership 0.449. Situation report drafted at docs/model/situation_report_20260918.md	output/model/sandbox/switches_root/s1_concave_benefit/	Torch rows to add at 09:00

Date	Label	What was done	Pointer	Status
2026-09-18 (02:40)	all	ψ re-normalized single-switch table (completed fertility 2.1 in every row; baseline loss 2607, ψ 0.190): rental wedge 402 (ownership 0.73, recent-parent gap 0.28, rooms response 0.97); earnings penalty 1746 (ψ 0.315, first birth 27.1, rooms response 0.85, three-plus gap 0.15); mortgage block 2050 (ownership 0.35, p90/p50 5.0); BGM earnings 2100 (childless 0.14, p90/p50 2.9); E6b earnings 2325 (p90/p50 3.38 at target 3.52); $\kappa_H=0$ 2379; maturation 2464; credit line 2539 (first birth 25.85, nearest the target). Full table output/ model/sandbox/ switches_root/ psi_root_ comparison.md; concave benefit and the Torch runs (estate, all switches) pending	output/model/ sandbox/ switches_root/	Report at 10:00

Date	Label	What was done	Pointer	Status
2026-09-18 (03:30)	E1, P5, X2	Earnings, estates and all-switches at fixed ψ. One PSID decomposition (E6b, both components after-tax): fertility 1.92, first birth 27.2, wealth/earnings 5.54, old p90/p50 3.44 (target 3.52; baseline 4.12), ownership 0.478, loss 2341. Boar–Gorea–Midrigan (no permanent types): fertility 2.05, childless 0.154, p90/p50 2.92 (the tail collapses without types), ownership 0.499, mean rooms 6.09, loss 2096. Estate transfer to ages 45–65 with net-of-selling-cost valuation: fertility 1.92, wealth 5.12, p90/p50 4.08, loss 2599, solve 54 min (outer loop). All switches together at fixed ψ : fertility 0.92, loss 11917, uninformative for the same reason as the maturation law alone (cost changes at fixed ψ); the ψ -root twin is queued on Torch	output/model/sandbox/switches_fixed/	Reading in the morning report

Date	Label	What was done	Pointer	Status
2026-09-18 (01:50)	all	Torch congested (about 2,800 jobs pending; the smoke waited 85 min unscheduled), so Muse stopped at the gate with the batch staged and hash-verified (sandbox_psi_root_20260918a). Lead decision: submit the ten ψ -root jobs anyway (same validated pattern as Tuesday's smoke) and run the six cheaper ψ -root specs locally in a chain overnight as insurance	output/model/sandbox/torch_psi_root/jobs.json, output/model/sandbox/switches_root/chain.log	Results by morning depend on the queue

Date	Label	What was done	Pointer	Status
2026-09-18 (00:40)	P3, H1, H3, H6	Single switches at fixed parameters and fixed ψ (baseline: fertility 1.872, childless 0.239, first birth 26.96, ownership 0.459, rooms response 1.07, recent-parent gap 0.45, loss 2535). Earnings penalty 20 percent while children at home: fertility 1.14, childless 0.46, first birth 30.2, loss 7176 (a cost change; read only with ψ re-normalized, running on Torch). Mortgage block (origination-only collateral test, 11 percent amortization, no cash-out): fertility flat 1.876, ownership 0.459→0.350, loss 2233. Rental wedge replacing the cap and χ (intercept 0.02, slope 0.05 above 6 rooms, $\chi=1$, cap 11): ownership 0.459→0.737, fertility 1.79, first birth 27.4, rooms response 0.91, recent-parent gap 0.27, loss 804 (a third of the baseline; the wedge intercept is a strong ownership lever and would be calibrated down). Modest credit line (2 d 0.10): fertility	output/model/ sandbox/ switches_fixed/	ψ -root twins on Torch; estate, all-switches and earnings specs running locally

Date	Label	What was done	Pointer	Status
2026-09-18 (00:25)	F3	ACS profile with the parent-age law and ψ re-normalized: children at home by parent age matches the ACS at 22–34 (1.17 vs 1.18 at 30; 1.43 vs 1.44 at 34), is slightly above it at 38–42, lies between the ACS any-age and under-18 series at 46–58, and is zero from 62 where the constant law kept a tail. The share with any child at home is somewhat high at 22–42. Both faults of the constant law (40 percent too few children at 30–42; adult “dependents” after 58) are gone. This is the evidence F3 asked for	output/model/sandbox/maturation_root/profile/dependents_by_age.png	Recommend adopting the parent-age law; refit later

Date	Label	What was done	Pointer	Status
2026-09-18 (00:15)	E1, H1	Part 5 specs written by Muse (16 files, fixed- ψ and ψ -root twins): modest credit line $\lambda_d = 0.1875$ (three quarters of a quarter of mean earnings, Kaplan–Violante); earnings E6b (one PSID decomposition, both components after-tax scaled, 15 states); earnings BGM (Boar– Gorea–Midrigan post-tax AR(1) 0. 964/0.150, no permanent types, transitory omitted, 5 states); single-switch and all-switches specs. Two sandbox blockers found (list parsing in the loader; a stale guard on the penalty), fix delegated	code/model/ sandbox/specs/	Loader fix in progress; then the runs

Date	Label	What was done	Pointer	Status
2026-09-18 (00:05)	F3	Parent-age law with ψ re-normalized to 2.1 (nine evaluations, 13 min): ψ 0.167 against the baseline root's 0.190; loss 2464 against 2607; childlessness 0.210 (baseline 0.201), first-birth age 26.5 (26.1), first births 30+ 0.259 (0.234); rooms gap three-plus vs one-to-two 0.57 (0.24; target 0.35), first-birth rooms response 1.39 (1.12; target 0.72); ownership 0.449, mean rooms 5.69. Children staying longer raises the space moments, which a refit would trade against a lower h_P . ACS profile for this solution running	output/model/sandbox/maturation_root/	Profile pending; baseline twin solving

Date	Label	What was done	Pointer	Status
2026-09-17 (23:45)	P5	Part 4 built and receipted: estate_ receiver: ages_ 45_65 (estates paid as an equal transfer to households aged 45–65 through an outer damped fixed point around the stationary solve; balances to 1e-10) and bequest_net_of_selling_cost: true; default off, bitwise-off test, full suite 227 passed. Cost note: each estate iteration re-solves the equilibrium	output/model/ opencode_tasks/ switch_batch_ part4_20260918/	Part 5 specs next, then the runs

Date	Label	What was done	Pointer	Status
2026-09-17 (23:30)	P3, H1, H3	Switch batch parts 1–3 built by Muse and receipted (output/model/opencode_tasks/switch_batch_part3_20260917/): child_earnings_penalty: $[\theta, \tau_1, \tau_2, \tau_3]$ (working-age earnings scaled in the budget; pensions, payroll base and down-payment test on unpenalized earnings); mortgage_origination_only: true and mortgage_amortization: α (stayer floor $\max(b(1-\alpha), b)$ when in debt, collateral floor at origination only, rollover taper bypassed when on); rental_wedge_intercept/slope/knee (renter budget only, user cost untouched). All default-off, bitwise-off tests, full suite 220 passed. Part 4 (estate receiver) and Part 5 (earnings and credit-line specs) launched next	commits e1a162ea and following	Lead review of parts 1–2 diffs still to do line by line

Date	Label	What was done	Pointer	Status
2026-09-17 (night)	F3, P1	Bellman fix delivered by Muse and verified (exact second housing-stage solve in fertile ages; constant mode bitwise; 201 tests; +25% Bellman cost in fertile ages). Its receipt records that V falls when a child stays one more period at every birth-destination state: at the retained parameters children are net costs while at home, and fertility is carried by the taste scales and the ψ normalization. Consequence: the fixed-ψ run of the parent-age law (completed fertility 0.49, childlessness 0.83, output/ model/sandbox/maturation/) is not an evaluation of the law but of that fact; the informative comparison re-normalizes ψ to 2.1 (running). This sharpens P1: the benefit term is below the cost of a child at home for every household	output/model/opencode_tasks/maturation_switch_fix_20260917/	Re-run with ψ root in progress; batch of four switches launched

Date	Label	What was done	Pointer	Status
2026-09-17 (night)	F3	Lead line-by-line check of the switch: per-age hazard, binomial and newborn-exempt transition matrices, and the bitwise-off path are correct; the forward step blends standard and exempt rows by the newborn share within a cell (an approximation, acceptable and to be stated). Rejected: the Bellman corrected birth-destination values by a uniform cross-state mean of the exemption gain, which misprices the birth decision. Correction delegated to Muse: a second housing-stage solve with the exempt continuation in fertile ages. Fertility rows from the switch are provisional until then	output/model/ opencode_tasks/ maturation_ switch_fix_ 20260917/	Fix in progress

Date	Label	What was done	Pointer	Status
2026-09-17 (evening)	F3	Parent-age maturation switch built by Muse, default off, bitwise-nested, 6 new tests plus 52 neighbour tests pass; exemption implemented as m-d draws, no new state. Override keys child_ maturation_ mode: parent_ age, mu_young, a_ rise, a_full	code/model/ intergen_ eqscale_seq_ optimized/ tests/test_ child_ maturation_ switch.py	ACS-profile test next; lead line-by-line check pending

Date	Label	What was done	Pointer	Status
2026-09-17	H1, X3	Frictionless benchmark, steady state, fixed parameters and fixed ψ. No down payment ($\phi=1$): ownership 30–55 0.459→0.557, completed fertility 1.872→1.866, first-birth age unchanged. Unsecured credit line (five years' earnings, no age taper), down payment kept: ownership unchanged (0.450), completed fertility 1.872→1.948, first-birth age 26.96→25.79, wealth/earnings 5.18→4.43. Reading: at these parameters the timing mechanism runs through liquidity, not through the collateral requirement; the down payment governs ownership, not births. Both relaxed together trips a solver dead-mass invariant (does not solve). A sandbox bug was fixed on the way: a scalar ϕ override was silently reset to 0.80 by the package's <code>n_parity length</code> check	output/model/sandbox/frictionless/	Mechanism statement for the slides needs rewording; author to read

Date	Label	What was done	Pointer	Status
2026-09-17 (late)	all	Sandbox aligned to the recipe's own assembly: two bugs fixed (supply elasticity 1.75→0.63; child consumption floor 0.48→0). Baseline now matches ψ to 3e-6 and ownership within 1 point, but two wealth-distribution rows still differ (recent-parent gap 0.56 vs 0.14; old p90/p50 4.05 vs 4.50) because the recipe starts from a cluster-only checkpoint. Plumbing stopped here. Rule: cluster replay = exact levels; sandbox = quick direction-and-size tests	code/model/sandbox/README.md	Closed by decision

Date	Label	What was done	Pointer	Status
2026-09-17	all	Baseline is now solid. Codex replayed the archived initial-state recipe on Torch (job 17923835, 14 min): all 13 moments and 17 parameters match exactly, loss 179.2984242480; slides PDF and figure sources mapped. Receipt: output/model/paper_baseline_sep14/replay_20260917/README.md. The sandbox's own assembly still lands elsewhere; it is being aligned to the recipe's functions so spec comparisons run at deck level	replay receipt	Sandbox alignment in progress

Date	Label	What was done	Pointer	Status
2026-09-17	P1, F2	main reconciled with the paper baseline (commit 40cceb4, solver hash 2992412...). The four fixed- ψ steady states rerun on it are identical to every digit to the earlier runs (losses 2535.40 / 8981.09 / 1950.60 / 2376.63), so the earlier sandbox conclusions stand unchanged and the solver version never affected the sequential household problem. The remaining gap to the deck's levels is entirely in the original driver's parameter assembly, pending the Codex replay	output/model/sandbox/baseline_code/	Conclusions unchanged
2026-09-17	E1	Author tentatively leans toward removing permanent earnings types and considering persistent-plus-iid-transitory earnings risk; final decision pending	this section	Tentative—author decision pending

Date	Label	What was done	Pointer	Status
2026-09-17	all	Branch question settled by the Codex session: the paper's code is the tagged worktree tmp/paper_baseline_sep14/ (paper-baseline-2026-09-14, hashes checked). All sandbox specs now set package_root to it. The numerical replay of the original recipe, and the testing framework, are owned by that session; the sandbox gate becomes a comparison against that replay	tmp/paper_baseline_sep14/PAPER_BASELINE.md	Replay pending on their side

Date	Label	What was done	Pointer	Status
2026-09-17 (early)	all	Gate rerun on the production snapshot with the receipt-hash solver: the joint (price, ψ, transfer) root converges in 14 evaluations (33 s each) but to a different equilibrium: ψ 0.163 vs 0.149, mean rooms 6.05 vs 6.42, ownership 30–55 0.478 vs 0.540, recent-parent gap 0.46 vs 0.14; fertility rows match within 0.01. Solver version is ruled out; the remaining difference is in the parameter construction, which the sandbox re-implements. Next step: replay the fetched run_capped_beta.py/run_profile.py locally with the snapshot as source_root instead of re-implementing, and diff the resulting parameter object against the sandbox's	output/model/sandbox/baseline_joint/	Gate still open; sandbox valid for differences, not levels

Date	Label	What was done	Pointer	Status
2026-09-17 (early)	all	Production snapshot fetched (367 .py files, corrected_initial_source_fetched/, MANIFEST with hashes): every package file matches commit 70abd4a8 except solver.py, of which three versions exist (main 7a0baaa2..., the template's corrected_initial_source 3bd6782e..., and the receipt's 2992412... found only in corrected_initial_source_v2, a folder the template's own provenance fields do not name). The receipt-hash solver is local as solver_receipt_hash.py. Gate rerun against the snapshot with that solver is running (joint root, full grid)	corrected_initial_source_fetched/ MANIFEST.md	Provenance fields in the template disagree with where the receipt hash lives; flag for the author

Date	Label	What was done	Pointer	Status
2026-09-16 (night)	all	main is not the production code. The retained initial state (job 17655042) ran a snapshot of branch codex/balanced-social-security at commit 70abd4a8 plus a corrected solver.py (sha256 2992412...); main's package differs from 70abd4a8 in 8 files (kernels.py, solver.py, parameters.py, utils.py; three modules absent on main), and neither solver hash matches the corrected one. Every sandbox run above therefore solved main's model, not the deck's; differences between specs remain informative, levels do not. The production snapshot is being fetched under output/model/e5f_final_night_20260913/corrected_initial_source_fetched/; the sandbox gets a package_root option to import it	this row	Author decision: which branch is the paper's code; merge or retarget main

Date	Label	What was done	Pointer	Status
2026-09-16 (evening)	all	Cluster back. Recovery stage-0 job 17858740 COMPLETED at its evaluation budget, candidate $\psi=0.1339$ unaccepted (production_ eligible: false, needs a continuation, author's call); queue empty. The retained initial state's driver was fetched from scratch into output/model/ e5f_final_ night_20260913/ corrected_ initial_ template_v6_ fetched/; its PAYGO helper e5f_stationary_ paygo.py is not on main but on branch codex/balanced- social-security (commit 1a5e88b7), while two main scripts import it	fetched folder MANIFEST.md	Merge decision for the author; gate rerun in progress
2026-09-16	F3	Children at home by parent age, model vs ACS: model 40% low at 30–42, tail above 58 mimics adult co-resident children	output/model/ sandbox/ dependents_by_ parent_age/	Memorylessness is not harmless; fix type is the author's call

Date	Label	What was done	Pointer	Status
2026-09-16	P1, F2	First sandbox comparison at fixed parameters: S2 halves fertility, S1 mild and improves loss, $\kappa_H = 0$ inert in the steady state (table below)	output/model/sandbox/	Read; author decision pending
2026-09-16	all	Sandbox for one-change GE stationary-state tests: make sandbox SPEC=name, four output files, switches for S1/S2 and κ_H ; about 15 min per full-grid solve locally	code/model/sandbox/README.md	Built; regression gate BLOCKED: the retained initial state was solved by a driver that exists only on cluster scratch (bind_initial_balanced_pension), fertility 2.002 vs 2.100 at the retained psi. Fetch that driver when the cluster is back. Until then compare specs against the sandbox's own baseline, not the deck.
2026-09-16	M40, M41	Guido's questions on transaction volumes and rooms per person over time added to the ledger	docs/model/POST_PRESENTATION_ISSUES.md	Open

Sandbox results, September 16 (fixed parameters, fixed ψ , one GE steady state per spec; levels are the sandbox's own baseline, not the deck's, see the gate note above).

Moment	Baseline	S2 SSK weighting	S1 log benefit	$\kappa_H = 0$
Completed fertility	1.872	1.016	1.920	1.871
Childless 40–44	0.239	0.486	0.244	0.239
Mean first-birth age	26.96	30.25	26.55	26.97

Moment	Baseline	S2 SSK weighting	S1 log benefit	$\kappa_H = 0$
First births at 30+	0.285	0.478	0.261	0.285
Ownership 30–55	0.459	0.474	0.453	0.452
Mean rooms	5.72	5.57	5.72	5.75
Rooms response, first birth	1.07	0.98	1.07	1.08
Rooms, 3+ vs 1–2 children	0.31	0.85	0.17	0.28
Recent-parent ownership gap	0.45	0.47	0.41	0.44
Old p90/p50	4.12	3.78	4.17	4.11
Wealth/earnings	5.18	5.20	5.21	5.18
Price	0.791	0.779	0.790	0.793
Loss (12 rows, sandbox)	2535	8981	1951	2377

Reading. **S2** (multiply utility by $e(m)$) more than doubles the cost of the first child at $\sigma = 2$: fertility halves, childlessness doubles, first births move four years later, and the three-plus families that remain are strongly selected (rooms gap 0.85). Direction as derived; magnitude means S2 is not a free correction, ψ would have to roughly double and the timing block re-estimated. **S1** ($\psi \log(1 + m)$) is mild: fewer large families (rooms gap 0.31 to 0.17), earlier first births, completed fertility up 0.05, loss down a fifth; the intensive margin now responds to preferences as intended. $\kappa_H = 0$ changes nothing in the stationary state (ownership down 0.6 points, loss slightly better): the tenure shock’s only footprint is numerical, which supports F2’s deterministic option; its policy-sign role must come from the transition, not the level. Outputs: output/model/sandbox/*_psi_fixed/ and the three compare_* folders.

Combined table

Costs and blocking flags are each reviewer’s own. “Same” means the two recommendations coincide in substance even when the one-word label differs.

Label	Fable	ChatGPT	Agreement	Cost (F / C)	Blocks recalibration (F / C)
P1 Utility form	Decide	Keep	Differ on action	Low / Low	Yes / Yes
P2 Cross-partial	Keep (restate on u)	Change (restate on u)	Same	None / Low	No / No
P3 First-birth cost	Decide, lean change	Keep	Differ	Low / Low	Yes / No

Label	Fable	ChatGPT	Agreement	Cost (F / C)	Blocks recalibration (F / C)
P4 Two child states	Keep	Keep	Same	None / Low	No / No
P5 Bequests	Change (net estate, receiver)	Keep (net estate)	Differ on receiver	Medium / Low	Yes / Yes
F1 Timing	Keep	Keep	Same	None / Low	No / No
F2 Taste shocks	Change ($\kappa_H = 0$ or estimate)	Change (fixed numerical)	Differ on replacement	Low–medium / Low	Yes / Yes
F3 Maturation	Change (parent-age hazard)	Change (child-stage state)	Differ on fix	Low–medium / Medium	Yes / Yes
F4 Orphans	Change via F3	Change (care pool if needed)	Same if F3 adopted	None / Medium	Yes / Yes
F5 Fecundity	Keep (decide unintended births)	Keep	Same	Low / Medium	No / Yes
H1 Constraint set	Change (amortization, origination LTV)	Change (description only)	Differ on depth	Medium / Medium	Yes / Yes
H2 Rollover	Change via H1	Change (contract rule)	Same direction	Low / Medium	Yes / Yes
H3 Rental cap	Change (wedge baseline)	Change (cap baseline, soft-tail robustness)	Differ on baseline	Medium / Medium	Yes / Yes
H4 Supply	Change (stock-flow)	Change (stock-flow)	Same	Medium / Medium	Transition / Yes
H5 Landlord	Keep and state	Change documentation	Same	Low / Medium	No / Yes
H6 Owner premium	Decide, lean wedge	Keep	Differ	Medium / Low	Yes / Yes
D1 Accounting	Keep and write	Change	Same task, differ on unit	Low / Medium	No / Yes
D2 Closure	Decide	Change	Complementary	Low or high / High	Transition / Yes
D3 Fiscal	Keep	Keep	Same	None / Low	No / Yes
D4 Geography	Change (national)	Change (national)	Same	Medium / Medium	Yes / Yes
E1 Earnings	Change (one decomposition)	Keep (reconcile first)	Same procedure	Low / Medium	Yes / Yes
E2 Risk and fertility	Keep	Keep	Same	None / Low	No / No

Label	Fable	ChatGPT	Agreement	Cost (F / C)	Blocks recalibration (F / C)
X1 Interactions	order: F3, E1; H1, F2; H4, D2; D4	order: units, closure; preferences, credit; earnings, scope	Differ on order	n/a	n/a
X2 Minimum set	7 changes, 8 explain, 6 decisions	package above; extensions conditional	Large overlap	n/a	n/a
X3 Missing	time cost, formation, widowhood, conventions	welfare population, omitted child costs, conventions	Disjoint, both included	n/a	n/a

Appendix. Verification status of claims made by one reviewer only

Fable’s code claims carry file and line in its review. ChatGPT’s claims were checked here where a local source exists:

- ChatGPT: the objective in Scholz, Seshadri and Khitatrakun contains $e U(c/e)$, the scale multiplying utility. Verified in docs/reference/scholz_seshadri_khitatrakun_2006_jpe.pdf, p. 615: “Expected lifetime utility is then $E[\sum \beta^{j-S} n_j U(c_j/n_j)]$ ”; p. 619: $n_j = (A_j + 0.7K_j)^{0.7}$.
- ChatGPT: Greaney et al. use an owner-service premium. Verified in the Zotero PDF (MAXJ699L): services $\chi h + h^r$, χ “captures the non-pecuniary benefits of ownership”, calibrated $\chi = 1.0506$ (pp. 10–11, Table 1).
- ChatGPT: Baum-Snow and Han report about 0.5 for floor space and 0.3 for units. Verified in the Zotero PDF (R5GI43BX), abstract; tract-level 0.42 and 0.35 on pp. 1898–1899. Correction to both reviews: a metro-aggregate floor-space elasticity of 0.61–0.63 appears on p. 1937, so the project’s 0.63 has a page, though not the headline object.
- ChatGPT: Kaplan, Mitman and Violante distinguish origination limits from limits on existing mortgages. Verified, p. 3295: after origination there is no requirement that the outstanding principal stay below λ_m times the home’s value; p. 3290: constraints bind only at origination.
- ChatGPT: Sommer–Sullivan–Verbrugge model household landlords. Not on disk; cited from memory by both.
- Fable: Venti and Wise (2004) on old-age housing equity; Henderson and Ioannides (1983) on the rental externality. Not on disk; from memory.